

Brandon London

Due date: 9/17/2019

Total points possible: 65

The goal of this project is to reinforce what is going on in the call stack, as well as refreshing ourselves on the unix environment and using a debugger to examine memory locations. All tasks should be done on the student unix server `delmar.umsl.edu`. If you have never used a debugger to examine raw memory locations, now is the time!

Task 1

Your task here is to examine the size of the activation record for a function. In particular, I want to know the size in bytes on the stack that is taken up by all the data that is not simply local variables being stored. To do this, write a C++ program that contains several functions that call each other in a sequence. These functions should take in at least one variable by value. It would also be useful to set some local variables to some specific values inside of them. Then using GDB or some other debugger, set a breakpoint at the end of the lowest-level function and examine the memory locations of the stack (probably using `x` in `gdb`). You should be able to get an understanding of how much space is taken up by non-variable data. Note that you should keep these functions fairly simple as it will make your life easier. For demonstrating your work, I want you to write up a justification for your answer, including relevant evidence. This should include screen captures of at least some display by the debugger of memory, along with a description of its interpretation.

```

#include <iostream>
int multiplybytwo(int);
int subtractthree(int);

int main() {
    int x = 1;
    int y = 2;
    int z = 3;

    std::cout<<multiplybytwo(z)<<std::endl;

    return 0;
}

int multiplybytwo(int r) {
    int x = r * 2;
    int y = subtractthree(x);
    return y;
}

int subtractthree (int e) {
    int x = e-3;
    return x;
}

```

```

[bo15zb@delmar 3780]$ uname -a
Linux delmar.ums1.edu 3.10.0-957.21.3.el7.x86_64 #1 SMP Tue Jun 18 16:35:19 UTC 2019 x86_64 x86_64 x86_64 GNU/Linux

```

The first step should be that we need to check the amount of Bytes in the registers This can be accomplished by the command **uname -a**. From the picture above we can see that this is a 86_64 machine that will contain 8 Byte registers. Next we can use gdb to look into the stack of my task1.cpp program. We can get the stack from the function **disas main**.

```

(gdb) disas main
Dump of assembler code for function main:
   0x000000000040078d <+0>:    push    rbp
   0x000000000040078e <+1>:    mov     rbp, rsp
   0x0000000000400791 <+4>:    sub     rsp, 0x10
   0x0000000000400795 <+8>:    mov     DWORD PTR [rbp-0x4], 0x1
   0x000000000040079c <+15>:   mov     DWORD PTR [rbp-0x8], 0x2
   0x00000000004007a3 <+22>:   mov     DWORD PTR [rbp-0xc], 0x3
   0x00000000004007aa <+29>:   mov     eax, DWORD PTR [rbp-0xc]
   0x00000000004007ad <+32>:   mov     edi, eax
   0x00000000004007af <+34>:   call    0x4007d4 <_Z13multiplybytwei>
   0x00000000004007b4 <+39>:   mov     esi, eax
   0x00000000004007b6 <+41>:   mov     edi, 0x601060
   0x00000000004007bb <+46>:   call    0x400620 <_ZNSolsEi@plt>
   0x00000000004007c0 <+51>:   mov     esi, 0x400680
   0x00000000004007c5 <+56>:   mov     rdi, rax
   0x00000000004007c8 <+59>:   call    0x400670 <_ZNSolsEPFRSoS_E@plt>
   0x00000000004007cd <+64>:   mov     eax, 0x0
   0x00000000004007d2 <+69>:   leave
   0x00000000004007d3 <+70>:   ret
End of assembler dump.
(gdb) break *main + 34
Breakpoint 1 at 0x4007af
(gdb) info b
Num      Type           Disp Enb Address            What
1        breakpoint      keep y   0x00000000004007af <main+34>
(gdb) r
Starting program: /home/bcl5zb/3780/a.out

```

As we can see the first thing to be pushed is the rbp which is the base pointer. As you can see the rsp is being put into the rbp holding the memory address of the current position of the stack. We have to step into each function until we are at the lowest function in the stack.

```
(gdb) disas multiplybytwo(int)
```

```
Dump of assembler code for function _Z13multiplybytwoi:
```

```
0x00000000004007d4 <+0>:    push    rbp
0x00000000004007d5 <+1>:    mov     rbp, rsp
0x00000000004007d8 <+4>:    sub     rsp, 0x20
0x00000000004007dc <+8>:    mov     DWORD PTR [rbp-0x14], edi
0x00000000004007df <+11>:   mov     eax, DWORD PTR [rbp-0x14]
0x00000000004007e2 <+14>:   add     eax, eax
0x00000000004007e4 <+16>:   mov     DWORD PTR [rbp-0x4], eax
0x00000000004007e7 <+19>:   mov     eax, DWORD PTR [rbp-0x4]
=> 0x00000000004007ea <+22>:   mov     edi, eax
0x00000000004007ec <+24>:   call    0x4007f9 <_Z13subtractthreei>
0x00000000004007f1 <+29>:   mov     DWORD PTR [rbp-0x8], eax
0x00000000004007f4 <+32>:   mov     eax, DWORD PTR [rbp-0x8]
0x00000000004007f7 <+35>:   leave
0x00000000004007f8 <+36>:   ret
```

```
End of assembler dump.
```

```
(gdb) nexti
```

```
0x00000000004007ec in multiplybytwo(int) ()
```

```
(gdb) disas multiplybytwo(int)
```

```
Dump of assembler code for function _Z13multiplybytwoi:
```

```
0x00000000004007d4 <+0>:    push    rbp
0x00000000004007d5 <+1>:    mov     rbp, rsp
0x00000000004007d8 <+4>:    sub     rsp, 0x20
0x00000000004007dc <+8>:    mov     DWORD PTR [rbp-0x14], edi
0x00000000004007df <+11>:   mov     eax, DWORD PTR [rbp-0x14]
0x00000000004007e2 <+14>:   add     eax, eax
0x00000000004007e4 <+16>:   mov     DWORD PTR [rbp-0x4], eax
0x00000000004007e7 <+19>:   mov     eax, DWORD PTR [rbp-0x4]
0x00000000004007ea <+22>:   mov     edi, eax
=> 0x00000000004007ec <+24>:   call    0x4007f9 <_Z13subtractthreei>
0x00000000004007f1 <+29>:   mov     DWORD PTR [rbp-0x8], eax
0x00000000004007f4 <+32>:   mov     eax, DWORD PTR [rbp-0x8]
0x00000000004007f7 <+35>:   leave
0x00000000004007f8 <+36>:   ret
```

```
End of assembler dump.
```

```
(gdb) stepi
```

```
0x00000000004007f9 in subtractthree(int) ()
```

```
(gdb) disas sub
```

```
sub_epsilon_src_nodes  sub_magnitudes      subtractthree(int)
```

```
(gdb) disas subtractthree(int)
```

```
Dump of assembler code for function _Z13subtractthreei:
```

```
=> 0x00000000004007f9 <+0>:    push    rbp
0x00000000004007fa <+1>:    mov     rbp, rsp
0x00000000004007fd <+4>:    mov     DWORD PTR [rbp-0x14], edi
0x0000000000400800 <+7>:    mov     eax, DWORD PTR [rbp-0x14]
0x0000000000400803 <+10>:   sub     eax, 0x3
0x0000000000400806 <+13>:   mov     DWORD PTR [rbp-0x4], eax
0x0000000000400809 <+16>:   mov     eax, DWORD PTR [rbp-0x4]
0x000000000040080c <+19>:   pop     rbp
0x000000000040080d <+20>:   ret
```

```
End of assembler dump.
```

In the above screenshot I use the debugger to step into each of the functions until I reach the lowest function on the stack, in this case it is the function subtract three and I keep stepping into subtract three until I reach +16. Once we are there we can look into the next 40 elements in hexadecimal, half word using the command **x/40xw \$rsp**. We can see that some function had to call main. Based on the program we can see that 1, 2, 3 were called early in the main. We are at the lowest called function and as we keep an eye and know the addresses., we believe that 0x7fffffff500 was one of the first things pushed on the stack in main. Now that we are in the lowest function the address 0x7fffffff4b0 is where the stack pointer is currently. Then we subtract these two to find the difference that the memory takes up. When we subtract those two memory addresses we get 80 Bytes which is the size of the stack for our program.

```
(gdb) x/40xw $rsp
0x7fffffff4b0: 0xffffffffe4e0    0x000007fff      0x004007f1      0x00000000
0x7fffffff4c0: 0x000000002      0x000000000      0x004008ad      0x00000003
0x7fffffff4d0: 0x000000000      0x000000000      0x00000000      0x00000006
0x7fffffff4e0: 0xffffffffe500    0x000007fff      0x004007b4      0x00000000
0x7fffffff4f0: 0xffffffffe5e0    0x000000003      0x000000002      0x00000001
0x7fffffff500: 0x000000000      0x000000000      0xf72113d5      0x000007fff
0x7fffffff510: 0x000000000      0x000000000      0xffffffffe5e8    0x000007fff
0x7fffffff520: 0x000000000      0x000000001      0x0040078d      0x00000000
0x7fffffff530: 0x000000000      0x000000000      0x5f205857      0xdc0755e0
0x7fffffff540: 0x004006a0      0x000000000      0xffffffffe5e0    0x000007fff
```

Code: #include <iostream>

int multiplybytwo(int);

int subtractthree(int);

int main() {

int x = 1;

int y = 2;

int z = 3;

std::cout<<multiplybytwo(z)<<std::endl;

return 0;

}

int multiplybytwo(int r) {

int x = r * 2;

int y = subtractthree(x);

return y;

}

int subtractthree (int e) {

int x = e-3;

```
    return x;
}
```

Task 2:

Write another program (in C++) that will allocate a local static array of integers and then a dynamic array of integers. Are they stored next to each other? You can examine this by examining the memory addresses where they are located. As described in class, on some systems the size of a dynamic array is actually stored in the bytes previous to a dynamically allocated array. Through some experiments on your own, try to see if this is true on delmar. Is this true or not true also for the local array? As in the first part, describe the procedure you used to test for this.

```
#include <iostream>

int main (){
    int array[5];
    int *darray = new int[5];

    return 0;
}
```

[Read 10 lines]

^G Get Help	^O WriteOut	^R Read File	^Y Prev Page	^K Cut Text	^C Cur Pos
^X Exit	^J Justify	^W Where Is	^V Next Page	^U UnCut Text	^T To Spell

```

[[bcl5zb@delmar 3780]$ ls
a.out task1.cpp task2.cpp task3.cpp
[[bcl5zb@delmar 3780]$ task2.cpp
-bash: task2.cpp: command not found
[[bcl5zb@delmar 3780]$ nano task2.cpp
[[bcl5zb@delmar 3780]$ disas main
-bash: disas: command not found
[[bcl5zb@delmar 3780]$ g++ task2.cpp
[[bcl5zb@delmar 3780]$ gdb -q a.out
Reading symbols from /home/bcl5zb/3780/a.out...(no debugging symbols found)...done.
(gdb) disas main
Dump of assembler code for function main:
    0x00000000040067d <+0>:    push    %rbp
    0x00000000040067e <+1>:    mov     %rsp,%rbp
    0x000000000400681 <+4>:    sub     $0x20,%rsp
    0x000000000400685 <+8>:    mov     $0x14,%edi
    0x00000000040068a <+13>:   callq   0x400530 <_Znam@plt>
    0x00000000040068f <+18>:   mov     %rax,-0x8(%rbp)
    0x000000000400693 <+22>:   mov     $0x0,%eax
    0x000000000400698 <+27>:   leaveq
    0x000000000400699 <+28>:   retq
End of assembler dump.
(gdb) █

```

The first step is again to start up gdb. Once the gdb is started we can disassemble the main function, Note we do not have any function other than main unlike the last question.

My program only has a few lines of disassembly code. Line one is the base pointer and line two is the stack pointer. Those are going to be the same for every program as they are the configuration for the stack. Line three is going to subtract 32 from our stack pointer, this is for the local static array. I figured you could use the command **b *main +4** and **print \$rsp** this will give us the beginning of the local array which

would be **0xffffffffe3a0** as you can see below.

```
Breakpoint 1 at 0x400681
[(gdb) print $rsp
No registers.
(gdb) r
Starting program: /home/bcl5zb/3780/a.out

Breakpoint 1, 0x00000000400681 in main ()
Missing separate debuginfos, use: debuginfo-install glibc-2.17-260.el7_6.5.x86_64
4 libgcc-4.8.5-36.el7_6.2.x86_64 libstdc++-4.8.5-36.el7_6.2.x86_64
[(gdb) print $rsp
$1 = (void *) 0x7ffffffe3a0
(gdb) disas main
Dump of assembler code for function main:
    0x0000000040067d <+0>:    push    %rbp
    0x0000000040067e <+1>:    mov     %rsp,%rbp
=> 0x00000000400681 <+4>:    sub     $0x20,%rsp
    0x00000000400685 <+8>:    mov     $0x14,%edi
    0x0000000040068a <+13>:   callq   0x400530 <_Znam@plt>
    0x0000000040068f <+18>:   mov     %rax,-0x8(%rbp)
    0x00000000400693 <+22>:   mov     $0x0,%eax
    0x00000000400698 <+27>:   leaveq  0
    0x00000000400699 <+28>:   retq
End of assembler dump.
(gdb) █
```

The dynamic array is located in the heap, in line 4 of our disassembly code, it moves 20 into the EDI register and then calls a function located at **0x400530**. This is malloc however in this machine it is called **_Znam@plt**, it has to be that because we are not calling anything else. EDI gets the value 20 because the number of bytes passed to new array. In the program we requested for 5 int so that would be 20 in decimal or **0x14** in hex. If we set another breakpoint on line 5 we can return the address that the array gave to our pointer. This is where the dynamic array is now on the heap.


```

[(gdb) print $rax
$2 = 4195965
[(gdb) b *main+13
Breakpoint 3 at 0x40068a
[(gdb) step
Single stepping until exit from function main,
which has no line number information.

Breakpoint 3, 0x00000000040068a in main ()
[(gdb) disas main
Dump of assembler code for function main:
    0x00000000040067d <+0>:    push    %rbp
    0x00000000040067e <+1>:    mov     %rsp,%rbp
    0x000000000400681 <+4>:    sub     $0x20,%rsp
    0x000000000400685 <+8>:    mov     $0x14,%edi
=> 0x00000000040068a <+13>:   callq   0x400530 <_Znam@plt>
    0x00000000040068f <+18>:   mov     %rax,-0x8(%rbp)
    0x000000000400693 <+22>:   mov     $0x0,%eax
    0x000000000400698 <+27>:   leaveq
    0x000000000400699 <+28>:   retq
End of assembler dump.
[(gdb) print $rsp
$3 = (void *) 0x7fffffffe380
(gdb) █

```

These two arrays would obviously will not be next to each other. The static array address is at **0xffffffffe3a0** and the dynamic array array is at **0x7fffffffe380**. I could not find anything on the size of the dynamic array when examining the memory before the array.

Code:

```

[bcl5zb@delmar 3780]$ ls
a.out task1.cpp task2.cpp task3.cpp
[bcl5zb@delmar 3780]$ vi task2.cpp
#include <iostream>

```

```

int main (){

    int array[5];
    int *darray = new int[5];

    return 0;
}

```

Task 3:

Write a program that prompts the user for two numbers and stores them in signed integers. The program should then add those two numbers together and store the result in a signed integer and display the result. Your program should then multiply them by each other and store the result in another b integer and display the result. Then do the same but with dividing the first number by the second. Display an error message to the screen if an operation has happened that does not result in a correct calculation. In other words, make sure to test your code for error cases. You can safely assume I will only give your program integers (I will give your program only decimal digits).

```
Starting program: /home/bcl5zb/3780/a.out
Please enter the first Integer: -2000000
Please enter in the second Integer:20000000
The first added to the second is: 18000000
EXCEPTION WHEN MULTIPLIED
The first divided by the second is: 0
ALL ARITHMETIC COMPLETED SUCCESSFULLY
[Inferior 1 (process 21227) exited normally]
(gdb)
```

```
the program being debugged has been started
Start it from the beginning? (y or n) y
Starting program: /home/bcl5zb/3780/a.out
Please enter the first Integer: 2
Please enter in the second Integer:0
The first added to the second is: 2
The first multiplied by the second is: 0
EXCEPTION WHEN DIVIDED
EXCEPTION ERROR
ALL ARITHMETIC COMPLETED SUCCESSFULLY
```

```
(gdb) r
Starting program: /home/bcl5zb/3780/a.out
Please enter the first Integer: 222222222222222222222222
Overflow Error
[Inferior 1 (process 22064) exited normally]
(gdb) -lllllllllllllllllllll
Undefined command: "-lllllllllllllllllllll". Try "help".
(gdb) r
Starting program: /home/bcl5zb/3780/a.out
Please enter the first Integer: -lllllllllllllllllll
Overflow Error
[Inferior 1 (process 22080) exited normally]
(gdb) r
Starting program: /home/bcl5zb/3780/a.out
Please enter the first Integer: 222222222222222222222222
Overflow Error
[Inferior 1 (process 22084) exited normally]
(gdb) r
Starting program: /home/bcl5zb/3780/a.out
Please enter the first Integer: 1
Please enter in the second Integer: llllllllllllllllllllll
Overflow Error
[Inferior 1 (process 22096) exited normally]
(gdb) r
Starting program: /home/bcl5zb/3780/a.out
Please enter the first Integer: 2
Please enter in the second Integer:-222222222222222222222222
Overflow Error
[Inferior 1 (process 22116) exited normally]
(gdb)
```

```

#include <iostream>
#include <stdexcept>
#include <bits/stdc++.h>
#include <stdlib.h>
signed int add(int, int);
signed int divide(int, int);
signed int multiply(int, int);
int main() {
    signed int a, b, c, d, m = 0;
    std::cout << "Please enter the first Integer: ";
    std::cin >> a;

    if ((a <= INT_MIN) || (a >= INT_MAX)) {
        std::cout << "Overflow Error " <<std::endl;
        exit (EXIT_SUCCESS);
    }
    std::cout << "Please enter in the second Integer:";
    std::cin >> b;
    if ((b <= INT_MIN) || (b >= INT_MAX)) {
        std::cout << "Overflow Error " <<std::endl;
        exit (EXIT_SUCCESS);
    }

    //addition
    try{
        c = add(a, b);
        std::cout << " The first added to the second is: " << c << std::endl;
    } catch(std::runtime_error & e) {
        std::cout << "EXCEPTION WHEN ADDING" << std::endl << e.what();
    }
    // multiplication
    try{
        m = multiply(a, b);
        std::cout << "The first multiplied by the second is: " << m << std::endl;
    } catch (std::runtime_error & e) {
        std::cout << "EXCEPTION WHEN MULTIPLIED" << std::endl;
    }
    //divide
    try {
        d = divide(a, b);

```

```

        std::cout << "EXCEPTION WHEN MULTIPLIED" << std::endl;
    }
    //divide
    try {
        d = divide(a, b);
        std::cout<< "The first divided by the second is: " << d << std::endl;
    } catch (std::runtime_error & e) {
        std::cout << "EXCEPTION WHEN DIVIDED" << std::endl << e.what();
    }

    std::cout << "ALL ARITHMETIC COMPLETED SUCCESSFULLY" << std::endl;
}

signed int add(int a, int b){
    // overflow handler
    if (((b > 0) && (a > (INT_MAX - b))) || ((b < 0) && (a < (INT_MIN - b)))) {
        throw std::runtime_error("RESULTS IN OVERFLOW WHEN ADDED!\n");
    }
    else{
        return a + b;
    }
}

signed int multiply(int a, int b){
    if ((b > INT_MAX / a) || (b < INT_MIN / a)){
        throw std::runtime_error("EXCEPTION ERROR\n");
    }
    else {
        return a * b;
    }
}

signed int divide(int a, int b) {
    if ((b == 0) || ((a == INT_MIN && (b == -1)))) {
        throw std::runtime_error("EXCEPTION ERROR\n");
    }
    else {
        return a / b;
    }
}

```

```

[bcl5zb@delmar 3780]$ vi task3.cpp
#include <iostream>
#include <stdexcept>
#include <bits/stdc++.h>
#include <stdlib.h>
signed int add(int, int);
signed int divide(int, int);
signed int multiply(int, int);
int main() {
    signed int a, b, c, d, m = 0;
    std::cout << "Please enter the first Integer: ";
    std::cin >> a;

```

```

if ((a <= INT_MIN) || (a >= INT_MAX)) {
    std::cout << "Overflow Error " <<std::endl;
    exit (EXIT_SUCCESS);
}
std::cout << "Please enter in the second Integer:";
std::cin >> b;
if ((b <= INT_MIN) || (b >= INT_MAX)) {
    std::cout << "Overflow Error " <<std::endl;
    exit (EXIT_SUCCESS);
}

//addition
try{
    c = add(a, b);
    std::cout << " The first added to the second is: " << c << std::endl;
} catch(std::runtime_error & e) {
    std::cout << "EXCEPTION WHEN ADDING" << std::endl << e.what();
}
// multiplication
try{
    m = multiply(a, b);
    std::cout << "The first multiplied by the second is: " << m << std::endl;
} catch (std::runtime_error & e) {
    std::cout << "EXCEPTION WHEN MULTIPLIED" << std::endl;
}
//divide
try {
    d = divide(a, b);
    std::cout<< "The first divided by the second is: " << d << std::endl;
} catch (std::runtime_error & e) {
    std::cout << "EXCEPTION WHEN DIVIDED" << std::endl << e.what();
}

std::cout << "ALL ARITHMETIC COMPLETED SUCCESSFULLY" << std::endl;
}

signed int add(int a, int b){
    // overflow handler
    if (((b > 0) && (a > (INT_MAX - b))) || ((b < 0) && (a < (INT_MIN - b)))) {
        throw std:: runtime_error("RESULTS IN OVERFLOW WHEN ADDED!\n");
    }
    else{

```

```
    return a + b;
}
}
```

```
signed int multiply(int a, int b){
    if ((b > INT_MAX / a) || (b < INT_MIN / a)){
        throw std::runtime_error("EXCEPTION ERROR\n");
    }
    else {
        return a * b;
    }
}
```

```
signed int divide(int a, int b) {
    if ((b == 0) || ((a == INT_MIN && (b == -1)))) {
        throw std::runtime_error("EXCEPTION ERROR\n");
    }
    else {
        return a / b;
    }
}
```